

# Supervised Learning for Gen Z Mental Health Score Prediction from Social Media Usage

Pankaj Kumawat *MSc Data Science*  
*Univ. of Europe for Applied Sciences*  
 Potsdam 14469, Germany

pankaj@ue-germany.de Submitted to Raja Hashim Ali *Department of Business*  
*Univ. of Europe for Applied Sciences*  
 Potsdam 14469, Germany  
 hashim.ali@ue-germany.de

## Abstract

This paper uses supervised regression to predict Gen Z mental health score from social media usage, platform behavior, demographics, and sleep-related screen time. The dataset contains 1,000,000 records with 12 columns, and the experiments use a reproducible sample of 30,000 records. Seven research questions are studied through baseline testing, full model comparison, preprocessing analysis, feature importance, ranking stability, robustness testing, and practical model selection. ElasticNet achieved the best pure metric result with MAE of 0.777, RMSE of 0.971, and  $R^2$  of 0.582. Linear Regression was selected as the final recommended model because it produced almost identical accuracy and the highest weighted decision score of 98.248. The most important predictor was daily usage hours, followed by addiction level features. Overall, simple linear models were strong enough for this task, while complex ensemble models did not provide a clear practical benefit.

## Index Terms

supervised learning, social media analytics, mental health score, regression, Gen Z, feature importance, model robustness, digital wellbeing, scikit-learn

## I. INTRODUCTION

Social media is now a normal part of daily life for Gen Z. Users spend time on platforms for communication, entertainment, education, news, and content creation. These usage patterns create structured data that can be studied with supervised machine learning. In this project, the goal is to predict `mental_health_score` from measurable social media and demographic variables.

Mental health score prediction is not a medical diagnosis task in this report. It is treated as a data science problem where behavioral signals are used to estimate a continuous score. The prediction is challenging because wellbeing is affected by many personal and external factors that are not fully captured in a social media dataset. Still, usage intensity, night usage, platform choice, addiction level, and screen time before sleep can provide useful signals for pattern analysis.

This study addresses the problem through seven research questions. It starts with baseline regression models and then compares a wider set of supervised learning methods. It also examines preprocessing effects, feature importance, ranking stability, robustness under noisy or missing data, and a final weighted decision matrix for practical model recommendation.

The main contributions of this work are the following:

- A complete benchmark of supervised regression models for Gen Z mental health score prediction.
- Evidence that simple linear and regularized models perform as well as or better than more complex tree-based models on this dataset.
- Feature importance analysis showing that daily usage hours is the strongest predictor.
- A practical model selection framework that combines accuracy, reliability, interpretability, runtime, and simplicity.

The rest of this paper is organized as follows. Relevant background and related work are reviewed in Section II. The dataset is described in Section III. Section IV defines the regression problem, and Section V presents the research questions. The methodology is explained in Section VI, and the experimental setup is given in Section VII. Results and discussion are reported in Section VIII. Finally, Section IX concludes the paper.

## II. RELATED WORK AND BACKGROUND

### A. Social Media and Digital Wellbeing

Social media use has been widely studied in relation to adolescent and young adult wellbeing. Existing studies show that social media can be linked with stress, anxiety, sleep problems, social comparison, and overall mental health, but the relationship is not simple or universal [1]. The effect often depends on the user, the platform, the purpose of use, and the surrounding lifestyle.

Recent reviews also warn that digital media variables should not be interpreted as direct clinical evidence [2]. This is important for the present report because the model predicts a score from structured behavior data, not a diagnosis. Therefore, the aim is to understand patterns and compare models, not to make medical claims.

### B. Machine Learning for Tabular Prediction

Supervised regression is suitable when the target variable is continuous and labelled examples are available. Common approaches for tabular regression include Linear Regression, Ridge Regression, ElasticNet, Decision Tree, Random Forest, and Gradient Boosting. These models are widely available through the scikit-learn library, which supports preprocessing pipelines, cross-validation, and standard evaluation metrics [3].

Linear models are useful because they are fast and easy to explain. Regularized models such as Ridge Regression and ElasticNet can improve stability when features overlap or when small coefficient penalties are helpful [4]. Tree-based methods and ensemble models can capture nonlinear relationships, but they may not always improve results when the main signal is close to linear.

### C. What Predicts Mental Health Score

In this dataset, the expected predictors include social media usage intensity, platform behavior, night usage, addiction level, and sleep-related screen time. Usage intensity is important because daily exposure may reflect how much social media is part of the user's routine. Addiction level is also relevant because it summarizes a stronger attachment to platform use.

Demographic and platform variables may still be useful, but their effect can be weaker or context-dependent. For that reason, this report uses feature importance analysis instead of assuming which variables matter most. The final interpretation is based on the trained models and generated outputs.

## III. DATASET DESCRIPTION

### A. Source and Size

The Gen Z Social Media Usage Dataset is stored locally as `data/genz_social_media_usage_1Million.csv`. It contains 1,000,000 rows and 12 columns. Each row describes one user-level record with demographic, usage, platform, and wellbeing-related values. Key facts are summarized in Table I.

TABLE I  
DATASET AT A GLANCE

| Property                | Value                                        |
|-------------------------|----------------------------------------------|
| Name                    | Gen Z Social Media Usage Dataset             |
| Source                  | Local course CSV dataset                     |
| Rows                    | 1,000,000                                    |
| Columns                 | 12                                           |
| Target                  | <code>mental_health_score</code>             |
| Task                    | Supervised regression                        |
| Domain                  | Digital wellbeing and social media analytics |
| Rows used for modelling | 30,000                                       |

### B. Feature Groups

The dataset contains 11 input features after excluding the target. The features can be organized into four main groups. Demographic features include `age`, `gender`, and `country`. Usage intensity features include `daily_usage_hours`, `avg_session_minutes`, and `num_platforms_used`.

Platform behavior features include `primary_platform`, `purpose`, `night_usage`, and `addiction_level`. Sleep-related behavior is represented by `screen_time_before_sleep`. These variables are suitable for regression because they provide a mix of numeric and categorical information about user behavior.

### C. Target Variable

The target variable is `mental_health_score`. It is treated as a continuous value, so the task is regression rather than classification. The goal is to learn how the input variables map to the target score and then evaluate how well different supervised models perform.

## IV. PROBLEM STATEMENT

The task is formulated as a supervised regression problem. Each sample consists of a feature vector  $\mathbf{x} \in \mathbb{R}^d$  created from one social media usage record. The objective is to learn a function  $\hat{f}$  that predicts the mental health score:

$$\hat{y} = \hat{f}(\mathbf{x}) \approx y \quad (1)$$

where  $y$  is the observed `mental_health_score`. The model is trained on a labelled dataset  $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^N$  and evaluated using a held-out test split and cross-validation. Performance is measured using Mean Absolute Error (MAE), Root Mean Squared Error (RMSE), and the coefficient of determination ( $R^2$ ).

## V. RESEARCH QUESTIONS

This study is guided by seven research questions covering model performance, feature behavior, robustness, and practical deployment.

- RQ1** How effectively can baseline supervised learning models predict mental health score?
- RQ2** Which supervised learning model achieves the best predictive performance?
- RQ3** How do preprocessing strategies affect model performance?
- RQ4** Which input features contribute most to prediction?
- RQ5** How does model ranking change across MAE, RMSE, and  $R^2$ ?
- RQ6** How robust is the best model under cross-validation, noise, and missing-data conditions?
- RQ7** Which model provides the best practical balance of accuracy, interpretability, reliability, runtime, and deployment suitability?

## VI. METHODOLOGY

The overall workflow is illustrated in Fig. 1. It covers data loading, sampling, preprocessing, model training, evaluation, feature analysis, robustness testing, and final model recommendation.

### Supervised Learning Pipeline




Fig. 1. End-to-end supervised learning pipeline used in this study, covering dataset loading, preprocessing, model training, metric evaluation, robustness testing, and practical model selection.

### A. Data Loading and Sampling

The dataset was loaded from CSV format using Python. A reproducible random sample of 30,000 records was used for modelling so that the experiments could run quickly and produce the same outputs each time. A fixed random seed of 42 was used across sampling, splitting, and model training.

The target column was separated from the input features. Numeric and categorical columns were processed differently because the dataset contains both continuous values and text labels. This avoids forcing categorical variables into an incorrect numeric format.

### B. Preprocessing Strategies

Three preprocessing strategies were compared in RQ3:

- 1) **Numeric Only + Scale**: Numeric features were imputed and scaled.
- 2) **Impute + OneHot + Scale**: Numeric features were imputed and scaled, while categorical features were imputed and one-hot encoded.
- 3) **Impute + OneHot, No Scale**: Categorical features were encoded, but numeric scaling was not applied.

All preprocessing was implemented through scikit-learn pipeline objects. This helps avoid data leakage because transformations are fitted only on the training data before being applied to test data.

### C. Models

The project evaluated seven regression models. These models cover simple baselines, linear models, regularized models, and tree-based methods. Table II summarizes the models used in the experiments.

TABLE II  
MODELS USED IN THE STUDY

| Model             | Role in the Study                           |
|-------------------|---------------------------------------------|
| Dummy Mean        | Mean prediction baseline                    |
| Linear Regression | Simple interpretable regression model       |
| Ridge Regression  | Regularized linear model                    |
| ElasticNet        | Regularized linear model with mixed penalty |
| Decision Tree     | Simple nonlinear tree model                 |
| Random Forest     | Bagging ensemble of decision trees          |
| Gradient Boosting | Boosting ensemble model                     |

### D. Evaluation Metrics

All models were evaluated using MAE, RMSE, and  $R^2$ .

$$\text{MAE} = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i| \quad (2)$$

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2} \quad (3)$$

$$R^2 = 1 - \frac{\sum_i (y_i - \hat{y}_i)^2}{\sum_i (y_i - \bar{y})^2} \quad (4)$$

MAE gives the average absolute prediction error. RMSE gives more penalty to larger errors.  $R^2$  measures how much variance in the target is explained by the model.

### E. Feature Importance Analysis

Feature importance was used to identify which inputs contributed most to prediction. The main output ranked the transformed model features by importance. This helps explain the model beyond just reporting accuracy.

### F. Robustness Testing

The best-performing model was tested under different conditions. The clean holdout result was compared with 5-fold cross-validation, 15% numeric noise, and 15% missing values. This simulates situations where real-world data may be noisy or incomplete.

### G. Multi-Criteria Decision Framework

For the final recommendation, models were scored using accuracy, reliability, interpretability, runtime, and simplicity. These criteria were combined into a weighted decision score. This makes the final model choice more practical than selecting only the best metric value.

## VII. EXPERIMENTAL SETUP

All experiments were implemented in Python using scikit-learn, pandas, NumPy, matplotlib, and seaborn. The random seed was fixed at 42 to make the results reproducible. The code generated all tables and figures saved in the project folders.

TABLE III  
SOFTWARE CONFIGURATION

| Component     | Detail                |
|---------------|-----------------------|
| Language      | Python 3.x            |
| ML framework  | scikit-learn [3]      |
| Data handling | pandas, NumPy         |
| Visualization | matplotlib, seaborn   |
| Random seed   | 42                    |
| Sample size   | 30,000 rows           |
| Validation    | Holdout and 5-fold CV |

## VIII. RESULTS AND DISCUSSION

### A. RQ1 – Performance of the Baseline

The first experiment tested basic supervised models against a simple mean baseline. Table IV shows that Ridge Regression and Linear Regression performed almost the same, with MAE of 0.777, RMSE of 0.972, and  $R^2$  of 0.582. Decision Tree performed slightly worse, while Dummy Mean performed much worse.

TABLE IV  
RQ1 – BASELINE HOLD-OUT PERFORMANCE

| Model             | MAE   | RMSE  | $R^2$ | Train Time |
|-------------------|-------|-------|-------|------------|
| Ridge Regression  | 0.777 | 0.972 | 0.582 | 0.204      |
| Linear Regression | 0.777 | 0.972 | 0.582 | 0.249      |
| Decision Tree     | 0.797 | 0.998 | 0.559 | 0.421      |
| Dummy Mean        | 1.217 | 1.503 | 0.000 | 0.236      |

The result shows that the input features contain real predictive signal. A simple mean predictor cannot explain the variation in mental health score. The linear models already achieve strong performance, which suggests that the main signal is easy to capture with simple regression.

### B. RQ2 – Best Performing Model

RQ2 compared all supervised models. ElasticNet achieved the best pure metric result, with MAE of 0.777, RMSE of 0.971, and  $R^2$  of 0.582. Ridge Regression and Linear Regression were nearly identical. Gradient Boosting, Random Forest, and Decision Tree did not improve over the top linear models.

The main finding is that model complexity did not add much value. Gradient Boosting was close in accuracy but slower. Random Forest and Decision Tree were weaker. This means the final choice should consider interpretability and practical use, not only tiny metric differences.




Fig. 2. RQ1: Error metric comparison for the baseline models. Linear and Ridge Regression show lower error than Decision Tree and much lower error than Dummy Mean.

TABLE V  
RQ2 – FULL MODEL COMPARISON ON HOLD-OUT

| Model             | MAE   | RMSE  | $R^2$ | Time  |
|-------------------|-------|-------|-------|-------|
| ElasticNet        | 0.777 | 0.971 | 0.582 | 0.100 |
| Ridge Regression  | 0.777 | 0.972 | 0.582 | 0.097 |
| Linear Regression | 0.777 | 0.972 | 0.582 | 0.108 |
| Gradient Boosting | 0.779 | 0.973 | 0.581 | 6.848 |
| Random Forest     | 0.786 | 0.981 | 0.574 | 1.857 |
| Decision Tree     | 0.797 | 0.998 | 0.559 | 0.219 |
| Dummy Mean        | 1.217 | 1.503 | 0.000 | 0.095 |




Fig. 3. RQ2:  $R^2$  comparison for all models. ElasticNet, Ridge Regression, and Linear Regression form the strongest group, while Dummy Mean explains no useful variance.

TABLE VI  
RQ3 – EFFECT OF PREPROCESSING STRATEGIES

| Preprocessing Strategy    | MAE   | RMSE  | $R^2$ |
|---------------------------|-------|-------|-------|
| Numeric Only + Scale      | 0.777 | 0.971 | 0.582 |
| Impute + OneHot + Scale   | 0.777 | 0.972 | 0.582 |
| Impute + OneHot, No Scale | 0.777 | 0.972 | 0.582 |

#### C. RQ3 – Sensitivity to Preprocessing

RQ3 tested whether preprocessing choices changed the result. Table VI shows that all three preprocessing strategies produced almost identical rounded performance. Numeric Only + Scale had RMSE of 0.971, while the two one-hot encoded pipelines had RMSE of 0.972.

This shows that the model result is stable across preprocessing choices. Adding categorical variables through one-hot encoding did not visibly improve the rounded result. The model mainly depends on the strongest numeric and encoded usage features.

#### D. RQ4 – Feature Significance

RQ4 analyzed which variables were most important. The strongest feature was `daily_usage_hours`, with importance value 1.089. This value was much higher than the other predictors, showing that daily usage intensity is the main driver in the model output.

The next strongest predictors were addiction level features. This makes sense because addiction level summarizes the user's relationship with social media usage. Platform, country, number of platforms, and night usage had smaller rounded importance values in this experiment.

#### E. RQ5 – Stability of the Ranking Across Metrics

RQ5 checked whether the model ranking changed across MAE, RMSE, and  $R^2$ . Table VIII shows that ElasticNet ranked first across all three metrics. Ridge Regression and Linear Regression were tied closely, with average rank 1.33.




Fig. 4. RQ3: Comparison of  $R^2$  scores across preprocessing strategies. The differences are very small, showing that the result is not dependent on one narrow preprocessing setup.

TABLE VII  
RQ4 – TOP FEATURE IMPORTANCES

| Feature                  | Importance |
|--------------------------|------------|
| daily_usage_hours        | 1.089      |
| addiction_level_High     | 0.068      |
| addiction_level_Low      | 0.017      |
| primary_platform_YouTube | 0.005      |
| country_India            | 0.004      |
| num_platforms_used       | 0.002      |
| night_usage              | 0.001      |

TABLE VIII  
RQ5 – CROSS-METRIC RANKING STABILITY

| Model             | MAE Rank | RMSE Rank | $R^2$ Rank | Avg. Rank |
|-------------------|----------|-----------|------------|-----------|
| ElasticNet        | 1        | 1         | 1          | 1.00      |
| Ridge Regression  | 1        | 2         | 1          | 1.33      |
| Linear Regression | 1        | 2         | 1          | 1.33      |
| Gradient Boosting | 4        | 4         | 4          | 4.00      |
| Random Forest     | 5        | 5         | 5          | 5.00      |
| Decision Tree     | 6        | 6         | 6          | 6.00      |
| Dummy Mean        | 7        | 7         | 7          | 7.00      |






Fig. 5. RQ4: Feature importance values for the trained model output. Daily usage hours dominates the ranking, followed by addiction level features.

The ranking is mostly stable. The same top group appears across the metrics, and the same weaker models stay at the bottom. This supports the idea that final selection can focus on practical differences between the top models.

#### F. RQ6 – Robustness to Noisy and Missing Data

RQ6 tested the best model under different data conditions. The clean holdout result had RMSE of 0.971 and  $R^2$  of 0.582. The 5-fold cross-validation result was close, with RMSE of 0.976 and  $R^2$  of 0.575.

TABLE IX  
RQ6 – ROBUSTNESS RESULTS

| Scenario                | RMSE  | $R^2$ | $R^2$ Drop |
|-------------------------|-------|-------|------------|
| Clean holdout           | 0.971 | 0.582 | 0.000      |
| 5-fold cross-validation | 0.976 | 0.575 | 0.007      |
| 15% numeric noise       | 0.984 | 0.572 | 0.010      |
| 15% missing values      | 1.065 | 0.498 | 0.084      |

The model was stable under cross-validation and numeric noise. Missing values caused the largest drop, reducing  $R^2$  from 0.582 to 0.498. This shows that data completeness is important for this task.




Fig. 6. RQ5: Heatmap of model rankings across MAE, RMSE, and  $R^2$ . ElasticNet ranks first, while Ridge Regression and Linear Regression remain very close.




Fig. 7. RQ6: Robustness comparison across clean holdout, cross-validation, noisy input, and missing-value scenarios. Missing data causes the largest performance decrease.

### G. RQ7 – Practical Recommendation of Models

RQ7 used a weighted decision matrix to select the most practical model. Linear Regression had the highest final score of 98.248. Ridge Regression scored 96.794, and ElasticNet scored 96.309.

TABLE X  
RQ7 – WEIGHTED DECISION MATRIX

| Model             | Acc.    | Rel.   | Int. | Run.    | Simple | Score  |
|-------------------|---------|--------|------|---------|--------|--------|
| Linear Regression | 99.812  | 99.300 | 95   | 99.694  | 95     | 98.248 |
| Ridge Regression  | 99.812  | 99.300 | 90   | 100.000 | 90     | 96.794 |
| ElasticNet        | 100.000 | 99.300 | 90   | 99.661  | 85     | 96.309 |
| Decision Tree     | 94.925  | 98.400 | 70   | 98.527  | 70     | 88.683 |
| Random Forest     | 98.120  | 99.900 | 45   | 79.408  | 45     | 79.733 |
| Gradient Boosting | 99.624  | 99.400 | 40   | 0.000   | 45     | 67.248 |
| Dummy Mean        | 0.000   | 42.500 | 100  | 99.783  | 100    | 53.467 |

Linear Regression was selected because it is accurate, interpretable, simple, fast, and easy to deploy. ElasticNet is the best pure metric model, but its advantage is very small. For a practical student project and explainable workflow, Linear Regression is the better final recommendation.




Fig. 8. RQ7: Weighted decision score for practical model selection. Linear Regression achieves the highest overall score.

## IX. CONCLUSION

This study investigated supervised learning models for predicting Gen Z mental health score from social media usage data. The dataset contained 1,000,000 records and 12 columns, and the experiments used a reproducible 30,000-row modelling sample. The tested models included Dummy Mean, Linear Regression, Ridge Regression, ElasticNet, Decision Tree, Random Forest, and Gradient Boosting.

The results showed that ElasticNet had the best pure metric performance, with MAE of 0.777, RMSE of 0.971, and  $R^2$  of 0.582. However, Linear Regression achieved almost the same predictive performance and received the highest weighted decision score. This made it the final recommended model for the project.

Feature importance showed that `daily_usage_hours` was the strongest predictor, followed by addiction level features. Preprocessing choices had very little impact on rounded performance, and model rankings were stable across MAE, RMSE, and  $R^2$ . Robustness testing showed that the model handled cross-validation and numeric noise well, but missing values caused a larger performance drop.

Future work could test the model on a separate external dataset, add more behavioral variables, and evaluate fairness across demographic groups. More advanced explanation methods such as SHAP could also be used to make the predictions easier to understand.

#### ACKNOWLEDGEMENTS

I thank Prof. Raja Hashim Ali for supervising this assignment and providing the Overleaf conference template and assignment framework.

#### REFERENCES

- [1] B. Keles, N. McCrae, and A. Grealish, "A systematic review: the influence of social media on depression, anxiety and psychological distress in adolescents," *International Journal of Adolescence and Youth*, vol. 25, no. 1, pp. 79–93, 2020.
- [2] C. L. Odgers and M. R. Jensen, "Annual Research Review: Adolescent mental health in the digital age: facts, fears, and future directions," *Journal of Child Psychology and Psychiatry*, vol. 61, no. 3, pp. 336–348, 2020.
- [3] F. Pedregosa et al., "Scikit-learn: Machine learning in Python," *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011. [Online]. Available: <https://jmlr.org/papers/v12/pedregosa11a.html>
- [4] T. Hastie, R. Tibshirani, and J. Friedman, *The Elements of Statistical Learning*, 2nd ed. Springer, 2009.
- [5] G. James, D. Witten, T. Hastie, and R. Tibshirani, *An Introduction to Statistical Learning*, 2nd ed. Springer, 2021.
- [6] A. E. Hoerl and R. W. Kennard, "Ridge regression: biased estimation for nonorthogonal problems," *Technometrics*, vol. 12, no. 1, pp. 55–67, 1970.
- [7] H. Zou and T. Hastie, "Regularization and variable selection via the elastic net," *Journal of the Royal Statistical Society: Series B*, vol. 67, no. 2, pp. 301–320, 2005.
- [8] L. Breiman, "Random forests," *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001.